

Maximum-likelihood

Vincent Lee

1/18/2020

Workshop

negative log-likelihood normal function declaration

Negative log-likelihood will minimise it

```
Negative.LL.Normal<-function(mu.Sig.parameters,Sample.Vector){  
  n      <- length(Sample.Vector)  
  sigma  <- mu.Sig.parameters[2]  
  part_a <- (n/2)*log(2*pi*sigma)  
  part_b <- 1/(2*sigma)  
  y      <- (Sample.Vector-mu.Sig.parameters[1])^2  
  value  <- part_a + part_b *(sum (y))  
  return (value)  
}
```

```
dataPath <- "C:/Users/vincentlee/Desktop/Non_linear_models/Week2"  
Norm.Sample.Vector<-read.csv(file=paste(dataPath,"sample_for_optimization.csv",sep="/"),header=TRUE,sep=";",as.is=TRUE)  
head(Norm.Sample.Vector)
```

```
## [1] 12.6567272 10.8362391  4.1087307 10.9831141 10.7719944 -0.9756093
```

```
var(Norm.Sample.Vector)
```

```
## [1] 24.75509
```

Optimization of negative LL

```
Optimized.Negative.Log.Likelihood.optim<-optim(c(7,4),  
  Negative.LL.Normal,  
  Sample.Vector=Norm.Sample.Vector,  
  method="L-BFGS-B",  
  hessian=TRUE,  
  lower=c(-Inf,0),  
  control=list(trace=1))
```

```
## iter    10 value 3024.314081  
## final   value 3022.953774  
## converged
```

```
# 10 iteration
```

```
c(Optimized.Negative.Log.Likelihood.optim$par,  
  Optimized.Negative.Log.Likelihood.optim$value,  
  Optimized.Negative.Log.Likelihood.optim$counts,  
  Optimized.Negative.Log.Likelihood.optim$convergence)
```

```
##                               function  gradient  
##    10.22552    24.73032 3022.95377    17.00000    17.00000    0.00000
```

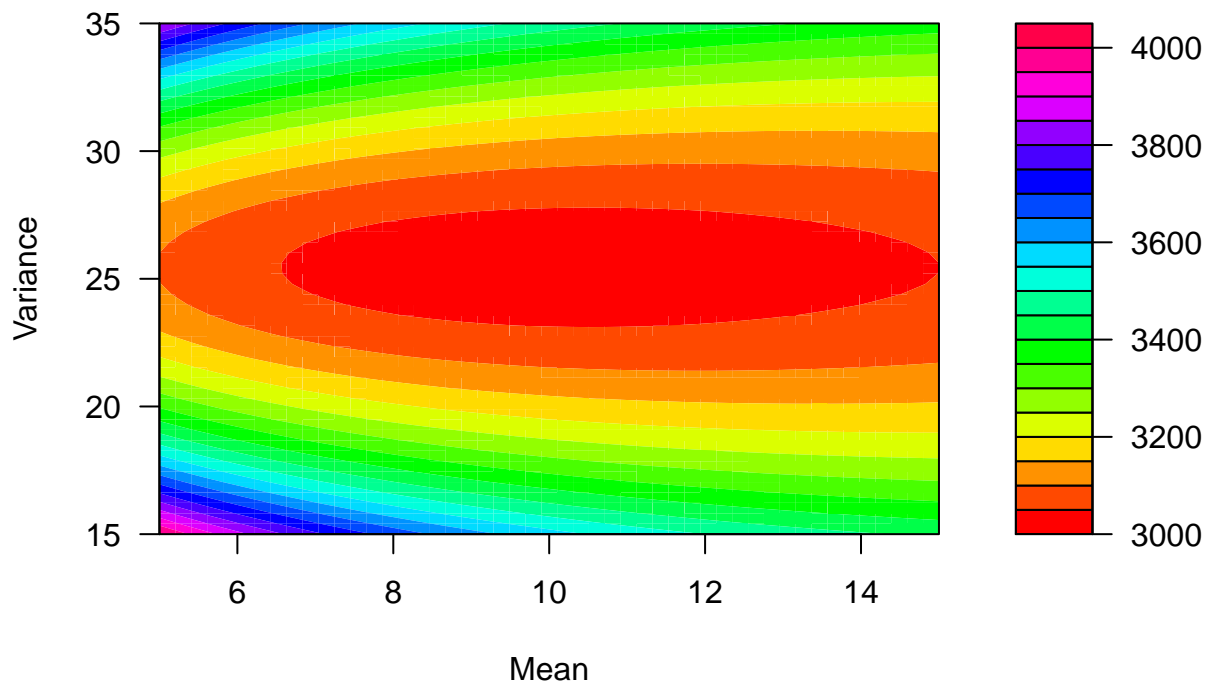
```
Optimized.Negative.Log.Likelihood.optim$hessian
```

```
##                [,1]                [,2]  
## [1,]  4.043619e+01 -1.136868e-07  
## [2,] -1.136868e-07  8.175431e-01
```

```
#Plot the objective function
```

```
Data.To.Plot.x<-seq(from=5,to=15,length.out=50)  
Data.To.Plot.y<-seq(from=15,to=35,length.out=50)
```

```
Term.1<-length(Norm.Sample.Vector)/2*outer(log(2*pi*Data.To.Plot.y),  
                                             rep(1,length(Data.To.Plot.y)))  
Term.2<-outer(rep(1,length(Data.To.Plot.x)),  
              unlist(lapply(Data.To.Plot.x,FUN=function(vector.element,Data.Vector)  
                          sum((Data.Vector-vector.element)^2),  
                          Norm.Sample.Vector)))*outer(1/2/Data.To.Plot.y,rep(1,length(Data.To.Plot.y)))  
  
Negative.Log.Likelihood.Data<-Term.1+Term.2  
filled.contour(Data.To.Plot.x,Data.To.Plot.y,Negative.Log.Likelihood.Data,  
               color.palette=rainbow,nlevels=20,xlab="Mean",ylab="Variance")
```



Analysis of the obtained estimates

Compare `Optimized.Negative.Log.Likelihood.optim$par` with `c(mean(Norm.Sample.Vector),var(Norm.Sample.Vector))`

```
rbind(Mean.Var=c(mean(Norm.Sample.Vector),var(Norm.Sample.Vector)),
      Optim.Output=Optimized.Negative.Log.Likelihood.optim$par)
```

```
##           [,1]      [,2]
## Mean.Var   10.22552 24.75509
## Optim.Output 10.22552 24.73032
```

Why `var(Norm.Sample.Vector)` is different from `Optimized.Negative.Log.Likelihood.optim$par[2]`?

```
n <- length(Norm.Sample.Vector)
rbind(Mean.Var=c(mean(Norm.Sample.Vector),var(Norm.Sample.Vector)),
      Optim.Output=c(Optimized.Negative.Log.Likelihood.optim$par[1],
                     Optimized.Negative.Log.Likelihood.optim$par[2]*(n/(n-1))))
```

```
##           [,1]      [,2]
## Mean.Var   10.22552 24.75509
## Optim.Output 10.22552 24.75508
```

Fisher score, fisher information

```
Biased.Var<- var(Norm.Sample.Vector)*((n-1)/n)
```

```
sum(Norm.Sample.Vector- mean(Norm.Sample.Vector))/var(Norm.Sample.Vector)
```

```
## [1] -2.583261e-14
```

```
(-n/(2*Biased.Var))+ (sum((Norm.Sample.Vector-mean(Norm.Sample.Vector))^2)/((Biased.Var^2)*2))
```

```
## [1] 3.552714e-15
```

Observed fisher's information

```
Optimized.Negative.Log.Likelihood.optim$hessian
```

```
##           [,1]      [,2]
## [1,]  4.043619e+01 -1.136868e-07
## [2,] -1.136868e-07  8.175431e-01
```

1. Element Hessian.1.1

```
n <- length(Norm.Sample.Vector)
n/var(Norm.Sample.Vector)
```

```
## [1] 40.39574
```

2. Element Hessian 1.2 and Element Hessian 2.1

```
sum(Norm.Sample.Vector- mean(Norm.Sample.Vector))/(var(Norm.Sample.Vector)^2)
```

```
## [1] -1.043527e-15
```

3. Element Hessian 2.2

```
Hessian2.2.1<- sum((Norm.Sample.Vector- mean(Norm.Sample.Vector))^2)/(Biased.Var^3)
(-n/(2*(Biased.Var^2)))+Hessian2.2.1
```

```
## [1] 0.8175421
```

Expected fisher's information (by assuming hessian $y-\mu = 0$)

```
#(Biased.Var<- var(Norm.Sample.Vector)*(length(Norm.Sample.Vector)-1)/length(Norm.Sample.Vector))
rbind(c(length(Norm.Sample.Vector)/Biased.Var,0),c(0,length(Norm.Sample.Vector)/2/Biased.Var^2))
```

```
##           [,1]      [,2]
## [1,] 40.43617 0.0000000
## [2,]  0.00000 0.8175421
```

4. Example simple linear regression

```
nSample<-500
sigmaEps<-1.5
set.seed(927436)
Eps<-rnorm(nSample,0,sigmaEps)
beta1<-1
beta0<-2.5
lambda<-0.5
X<-rexp(nSample,lambda)
Y<-beta0+beta1*X+Eps
# plot(X,Y)
```

```
dtf<-data.frame(X=X,Y=Y)
head(dtf)
```

```
##           X           Y
## 1 1.980506 4.0826325
## 2 2.542126 5.2707668
## 3 1.040133 6.9997564
## 4 1.694994 0.7666407
## 5 0.102001 1.9767673
## 6 2.751328 6.4738524
```

4.2 loglikelihood function

```
linModLL<-function(Parameters,regSample){
  # product of density of response f (y, theta)
  # L(theta, y) = sum (log f(yi, theta))
  # = sum( log(dnorm(sample, mean=bo+b1xi, sd=sigmaEps)))
  # Parameters[1] = beta0 Paramaters[2] = beta1 Sigma = Parameters[3]
  y <- regSample[,2]
  x <- regSample[,1]
  average <- Parameters[2]
  intercept <- Parameters[1]
  standev <-Parameters[3]
  log_dnorm <- log(dnorm(x = y, mean = intercept+(average*x),sd =standev))
  log_lik <- -sum(log_dnorm)
  return(log_lik)
}
```

```
linModLL(c(Beta0=beta0+1,Beta1=beta1+1,Sigma=sigmaEps),dtf)
```

```
## [1] 2523.769
```

```
linModLL(c(Beta0=beta0,Beta1=beta1,Sigma=sigmaEps),dtf)
```

```
## [1] 931.3951
```

4.3 FIT

```
Optimized.linModLL.optim <-optim(c(Beta0 = beta0,Beta1= beta1,Sigma=sigmaEps),  
                                linModLL,  
                                regSample=dtf,  
                                method="L-BFGS-B",  
                                hessian=TRUE,  
                                lower=c(-Inf,0.00),  
                                control=list(trace=1))
```

```
## final value 930.111159
```

```
## converged
```

```
# this is optim function
```

```
# we need to use newton approach for optimization in Assignment
```

```
Optimized.linModLL.optim$par
```

```
##      Beta0      Beta1      Sigma
```

```
## 2.4929634 0.9761973 1.5547015
```

Compare the results with linear model fit.

```
linM<-lm(Y~X,dtf)  
c(linM$coefficients,summary(linM)$sigma)
```

```
## (Intercept)          X
```

```
## 2.4929604 0.9761977 1.5578206
```

Assignment

```
my.Optimizer<-function(Start.Value,Function.To.Optimize,Epsilon,projectID){  
  iteration <- 10000 #random number  
  derivative <- function(Start.Value, Epsilon, projectID) {  
    # derivative function of testFunction  
    return((Function.To.Optimize((Start.Value + Epsilon),projectID) -  
      Function.To.Optimize((Start.Value - Epsilon), projectID)) / (2 * Epsilon))  
  }  
}
```

```

for (i in 1:iteration){
  update.step <- (Function.To.Optimize(Start.Value,projectID))/(derivative(Start.Value, Epsilon, projectID))
  #once update step is smaller than Eps, print results
  if (abs(update.step) < Epsilon){
    break
  }
  Start.Value<- Start.Value - update.step
}
root <- Start.Value
return (root)
}

```

To check your optimizer create a test function that needs to be optimized. In this project we use one-dimensional optimization, i.e. optimization with respect to only one variable. Add one more argument to the function, called projectID. The meaning of it will become clear in section Test. The function should cross x-axis at least in one point.

```

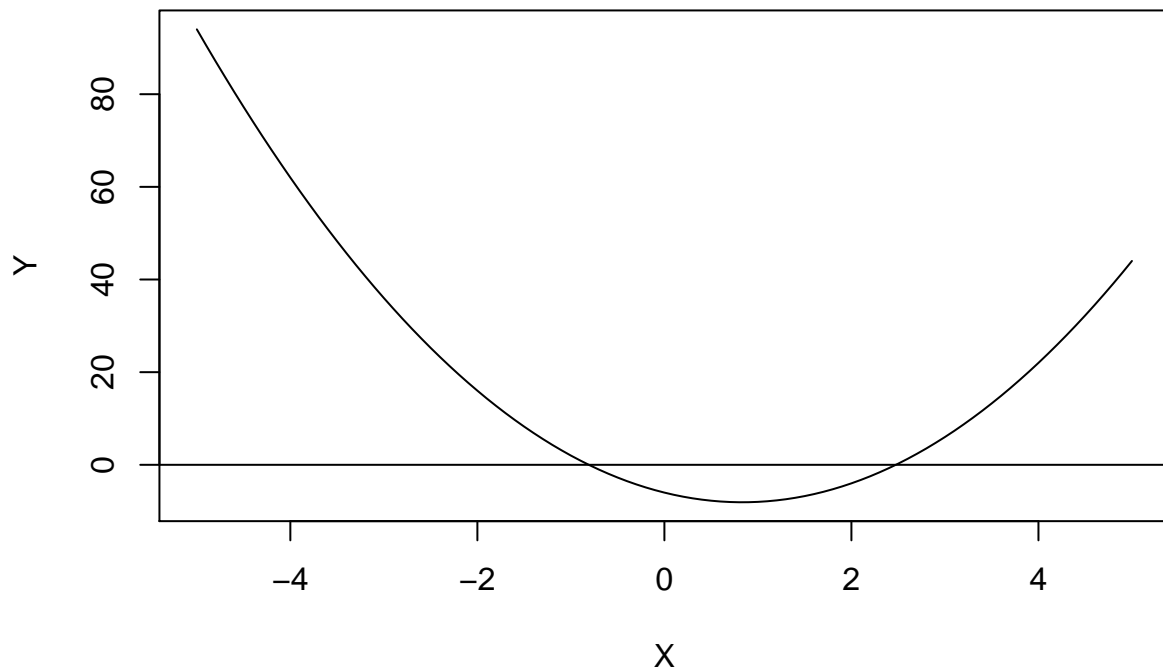
my.Function<-function(my.X,projectID) {
  my.X^2*3-my.X*5-6
}

```

```

X<-seq(from=-5,to=5,by=.1)
Y<-my.Function(X)
plot(X,Y,type="l")
abline(h=0)

```



You can also test the optimizer by running `uniroot()`.

```
uniroot(my.Function,lower=-5,upper=+1, tol=0.0001)
```

```
## $root
## [1] -0.8081429
##
## $f.root
## [1] -1.138463e-06
##
## $iter
## [1] 9
##
## $init.it
## [1] NA
##
## $estim.prec
## [1] 5e-05
```

```
my.Optimizer(-5, my.Function,0.0001,656)
```

```
## [1] -0.808143
```

Test

```
testFunction<-readRDS(file=paste(dataPath,"Week2_TestFunction.rds",sep="/"))$Week2_Test_Function
#project 656
testFunction(0, 656)
```

```
## [1] -5
```

```
#readRDS(file=paste(dataPath,"Week2_TestFunction.rds",sep="/"))
```

Make sure that declaration of your optimizer function contains `projectID` argument.

```
my.Optimizer(Start.Value=-100,
             Function.To.Optimize = testFunction,
             Epsilon=0.0001,
             projectID=656)
```

```
## [1] -5.000001
```

where:

`Start.Value` is initial guess for the optimizer, `testFunction` is the name of the test function that needs to be optimized, `Epsilon` is stopping criterion (set `Epsilon=0.0001`), `projectID` is individual project ID.

Find root (`my.Optimizer.root`) of the test function using your optimizer.


```
my.Optimizer(Start.Value=-100,
             Function.To.Optimize = testFunction,
             Epsilon=0.0001,
             projectID=656)
```

```
## [1] -5.000001
```

Find root (uniroot.root) using uniroot(). Use Epsilon=0.0001 as tolerance parameter (tol = 0.0001) of uniroot()

```
(uniroot.val<- uniroot(testFunction, c(-100, 0),tol = 0.0001, projectID=656))
```

```
## $root
## [1] -5
##
## $f.root
## [1] 1.007208e-07
##
## $iter
## [1] 14
##
## $init.it
## [1] NA
##
## $estim.prec
## [1] 5e-05
```

```
(my.optimizer.root<-optim(par=-100,
                          my.Optimizer,
                          Function.To.Optimize=testFunction,
                          Epsilon = 0.0001,
                          projectID = 656,
                          method="L-BFGS-B",
                          hessian=TRUE,
                          lower=c(-1e10,0),
                          control=list(trace=1)))
```

```
## final value -5.000001
## converged
```

```
## $par
## [1] -100
##
## $value
## [1] -5.000001
##
## $counts
## function gradient
##      2      2
##
## $convergence
```

```
## [1] 0
##
## $message
## [1] "CONVERGENCE: REL_REDUCTION_OF_F <= FACTR*EPSMCH"
##
## $hessian
##           [,1]
## [1,] -2.065015e-08
```

```
res <- list(Start.Value = -100,
            my.Optimizer.root = my.optimizer.root$value,
            uniroot.root = uniroot.val$root,
            uniroot.lower = -5,
            uniroot.upper = 0)
```

```
write.table(res, file = paste(dataPath, 'result.csv', sep = '/'), row.names = F)
```